

Searching for Task-Specific Vision Paths

Evolutionary Block Pruning Across Vision-Language Models

Tarun Tomar

<https://github.com/TarunTomar122/vision-pathways>

Abstract

Vision-language models normally execute the same complete vision encoder for every question, even when OCR, counting, object, attribute, and spatial queries may not require identical computation. We study whether fixed-budget combinations of vision blocks can be skipped without fine-tuning. A shared K-block route skips one searched set of exactly K blocks for every question, while a capability-specific K-block policy selects one same-size route using a known capability label. We introduce a source-balanced evolutionary search and compare it with independent ranking, contiguous removal, and random routes at matched budgets. Experiments use Qwen2.5-VL-3B-Instruct, SmolVLM2-2.2B-Instruct, and an 876-example image-disjoint selection split. Search transfers across architectures: on SmolVLM2, the searched shared four-block route beats independent construction by 4.91 percentage points. Capability specialization is less stable. On Qwen, the six-block capability policy beats the shared route by 2.17 points, driven by a 7.10-point OCR gain. On sealed IIIT5K, however, the SmolVLM2 OCR-specific route trails its shared route by 13.6 points. Combinatorial search reliably improves route construction, but capability labels do not define universally transferable vision pathways.

1 Introduction

Vision-language models (VLMs) have become a dominant paradigm for multimodal AI. Guided by scaling laws and advances in multimodal pretraining, VLMs have become increasingly powerful, but they have also become expensive to deploy [Kaplan et al., 2020, Hoffmann et al., 2022, Bai et al., 2025]. Unlike language-only models, a VLM must first convert an image into visual tokens and process them through a vision encoder before generating an answer. Every question normally executes the same complete stack of vision blocks, even when the visual information requested by the question is very different.

Capabilities like OCR, counting, and spatial reasoning may not require identical computation. This raises a simple question: does every capability need every vision block? If some vision blocks are redundant for a given input, replacing them with identity operations could reduce executed vision depth without introducing sparse kernels or changing the language decoder. However, a block that looks safe to skip by itself may not remain safe when several blocks are skipped together because residual blocks interact.

Existing efficiency work has explored quantization, structured layer pruning, and visual-token reduction. Layer-pruning methods such as ShortGPT and SliceGPT reduce depth or width in language models [Men et al., 2024, Ashkboos et al., 2024], while DynamicViT, Token Merging, SparseVLM, and VScan reduce the number of visual tokens processed by transformers [Rao et al., 2021, Bolya et al., 2023, Zhang et al., 2025b,a]. Short-LVLM directly studies training-free layer pruning in large vision-language models, but searches for one generally compressed model [Ma et al., 2025]. It therefore remains unclear whether the best removable *vision-encoder* blocks depend on the requested visual capability, and whether independently measured block importance composes when several blocks are skipped together.

We study two questions:

1. At the same number of skipped blocks, can combinatorial search find stronger routes than independent ranking, contiguous removal, or random selection?
2. Can capability-specific routes outperform one shared route across OCR, counting, object, attribute, and spatial questions?

A route is a set of exactly K vision-transformer blocks replaced by identity operations. A shared K-block route skips the same blocks for every question, while a capability-specific K-block policy uses one same-size route for each known capability label. We search for the best routes found under a frozen evolutionary procedure rather than claiming a global optimum. All comparisons use matched pruning budgets, meaning every compared route skips exactly the same number of blocks.




Figure 1: Accuracy drop after replacing one Qwen2.5-VL-3B vision block at a time with identity. Most blocks cause limited average damage when skipped alone, but sensitivity varies by capability and depth. This is screening evidence: a low single-block drop does not imply that several such blocks can be safely combined or that a capability is stored in a specific block.

We use Qwen2.5-VL-3B-Instruct as the primary model and SmolVLM2-2.2B-Instruct as a second architecture. Our dataset covers five visual capabilities: OCR, counting, object existence, attributes, and spatial relationships. We first measure the sensitivity of every vision block independently, then use source-balanced evolutionary search to construct shared and capability-specific routes at matched skip budgets. The evaluation separates route search, image-disjoint method selection, matched controls, and a post-freeze IIIT5K source-transfer test.

We find that evolutionary search consistently constructs stronger routes than naive block selection, but capability-specific routing is not universally better. On Qwen, the capability-specific six-block policy is 2.17 percentage points above the shared six-block route, driven by a 7.10-point OCR gain. On SmolVLM2, evolutionary search still helps: the searched shared four-block route beats independent construction by 4.91 points. However, its capability-specific policy is only 0.80 points better overall because counting and spatial gains are offset by a 13.55-point OCR loss. On sealed IIIT5K, the Smol OCR-specific route falls another 13.6 points below the shared route. These results suggest that route construction is genuinely combinatorial, but capability labels alone do not define stable or transferable vision pathways.

Our contributions are:

1. A source-balanced evolutionary procedure for searching fixed-budget vision-block routes.
2. A matched-budget study across two VLM architectures and five visual capabilities.
3. A capability-level, cross-model, and source-transfer analysis showing both the potential and limitations of capability-specific routing.

2 Motivation

2.1 Are individual vision blocks redundant?

Our first question was whether every vision block contributes equally to the model’s predictions. We tested this on Qwen2.5-VL-3B-Instruct, whose vision encoder contains 32 transformer blocks [Bai et al., 2025]. Starting from an unpruned baseline accuracy of 81.62% on 1,480 examples, we replaced each vision block with an identity operation, one block at a time, and measured the resulting accuracy drop. The benchmark contains questions from five capability buckets: OCR, counting, object existence, attributes, and spatial relationships. This produces a block-by-capability sensitivity map rather than a pruning policy: each intervention measures what happens when one block is removed while the other 31 remain intact.

The result is highly uneven across depth (Figure 1). Skipping block 28 reduces overall accuracy by only 0.14 percentage points, with no capability losing more than 1.67 points. In contrast, skipping block 15 reduces overall accuracy by 21.08 points. Blocks 0 and 31 are also sensitive, causing drops of 17.36 and 7.77 points respectively. The capability-level view is even less uniform. OCR accuracy falls by 25.59 points when block 14 is skipped and by 63.82 points when block 15 is skipped, while many other block-capability pairs show little measured damage.

These observations suggest that the vision encoder contains some removable computation, but they do not show that a capability is stored in a particular block. A block can look unimportant because nearby blocks compensate

for its absence, because the benchmark does not expose all of its functions, or because its contribution only becomes necessary after another block is removed. We therefore use the single-block sweep only as evidence of heterogeneous sensitivity and as a source of candidate blocks. It motivates pruning, but it cannot determine a safe multi-block route by itself.

2.2 Why independent block rankings do not compose

The most direct way to construct a deeper-pruned model is to rank blocks by their individual accuracy drops and combine the apparently safest ones. This follows the general intuition behind independent layer-importance methods such as ShortGPT [Men et al., 2024]. If block effects were additive, selecting the K least damaging blocks would be a reasonable approximation to the best K -block route.

Our initial progressive ablation shows why this assumption is unreliable. Skipping block 28 alone loses 0.14 points. The two-block route $\{5, 28\}$ loses 0.95 points, the three-block route $\{5, 9, 28\}$ loses 2.97 points, and the four-block route $\{3, 5, 9, 28\}$ loses 3.72 points. The final route removes 12.5% of the vision blocks and improves measured vision-encoder latency by 4.89%, but it loses 7.94 points on OCR. Blocks that appear safe under separate interventions can therefore become damaging when composed.

This failure is expected in a sequential residual network. Each block receives the representation produced by all preceding blocks, so removing one block changes the input distribution seen by every block after it. Let $\Delta(S)$ denote the accuracy loss after skipping a set of blocks S . For two blocks i and j , we can describe their pairwise interaction as

$$I(i, j) = \Delta(i, j) - \Delta(i) - \Delta(j).$$

A positive value means that the pair is more harmful than their independent effects predict, while a negative value means that the blocks are more compatible to skip together. Higher-order routes can contain interactions that no collection of one-block scores can recover. Block importance is therefore conditional on which other blocks are already skipped. The object we need to select is not a list of blocks, but a complete route.

2.3 Why route selection is a combinatorial problem

Once route quality depends on block combinations, evaluating blocks independently is insufficient. For a vision encoder with L blocks and a fixed skip budget K , the number of possible routes is

$$|\Omega_K| = \binom{L}{K}.$$

For Qwen’s 32-block vision encoder, this gives 35,960 possible four-block routes, 906,192 six-block routes, and 10,518,300 eight-block routes. Exhaustively running every route through a VLM benchmark would require millions of full inference evaluations at the larger budgets. The search space becomes even more expensive if a separate route is considered for each capability.

This motivates a fixed-budget combinatorial search. The method must compare routes with the same value of K , evaluate interactions among skipped blocks, and balance average performance against severe failures on individual data sources. It must also support two distinct hypotheses: one shared route may work well for all questions, or different capability labels may favor different routes of the same size. Finally, searched routes must be compared against independent rankings, contiguous removals, and random routes under exactly matched skip budgets. These requirements lead us to the source-balanced evolutionary search described next.

3 Methodology

3.1 Identity block skipping

Let a frozen vision encoder contain L transformer blocks $F_0, \dots, F_{L-1}$. We represent a route by a subset of block indices

$$S \subseteq \{0, \dots, L-1\}, \quad |S| = K,$$

where S contains the blocks to skip and K is the skip budget. For hidden state h_l , executing route S changes the vision encoder to

$$h_{l+1} = \begin{cases} h_l, & l \in S, \\ F_l(h_l), & l \notin S. \end{cases}$$

Selected blocks are therefore replaced by identity operations. We do not remove the patch embedding, token merger, language decoder, or any block parameters from the checkpoint, and we do not fine-tune the model. This intervention measures the effect of reducing executed vision depth while leaving the rest of the inference pipeline unchanged. Within each comparison, every route contains exactly K unique block indices.

3.2 Shared and capability-specific routes

We study two route families. A **shared route** uses one set S_{shared} for every question, regardless of its capability. This is the standard generic pruning setting: one fixed compressed pathway must preserve average behavior across the full benchmark.

A **capability-specific policy** contains one route S_t for each capability in the set

$$\mathcal{T} = \{\text{attribute, counting, object, OCR, spatial}\}.$$

At inference, the route associated with the known capability label is selected before the image passes through the vision encoder. Every S_t has the same skip budget K , so a policy cannot gain accuracy by executing more blocks for a difficult capability. This setup tests whether different capabilities prefer different fixed block subsets.

The policy is an oracle-label evaluation, not a learned router. Our experiments assume the capability label is already available and do not measure classification errors or routing overhead. The search also returns one frozen route per label rather than choosing a new route for every image or question.

3.3 Source-balanced objectives

Raw benchmark accuracy can hide source imbalance because some datasets contribute many more questions than others. We therefore evaluate each route against the full model on exactly paired examples and aggregate at the capability-source level. Let $I_{c,d}$ be the examples belonging to capability c and dataset source d . If $y_i(S)$ is one when route S answers example i correctly and zero otherwise, its cell accuracy is

$$A_{c,d}^S = \frac{1}{|I_{c,d}|} \sum_{i \in I_{c,d}} y_i(S).$$

The paired damage relative to the full model is measured in percentage points:

$$\Delta_{c,d}(S) = 100 (A_{c,d}^0 - A_{c,d}^S).$$

Positive values indicate lost accuracy, while negative values mean that the skipped model happened to outperform the full model in that cell. Every capability-source cell receives equal weight regardless of its number of examples. Failed or malformed model outputs count as incorrect.

For a shared route, let G be all capability-source cells and define its mean damage, worst-source damage, and source variability as

$$\mu(S) = \frac{1}{|G|} \sum_{g \in G} \Delta_g(S), \quad w(S) = \max_{g \in G} \Delta_g(S),$$

$$\sigma(S) = \sqrt{\frac{1}{|G|} \sum_{g \in G} (\Delta_g(S) - \mu(S))^2}.$$

Evolutionary survival minimizes the vector $q_{\text{shared}}(S) = (\mu, w, \sigma)$. When finalists must be ordered, we use the frozen scalar loss

$$L_{\text{shared}}(S) = 0.50\mu(S) + 0.30w(S) + 0.20\sigma(S).$$

Evolutionary search at a matched skip budget




Figure 2: Fixed-budget evolutionary route search. Candidate routes are evaluated with paired source-aware objectives, filtered by Pareto survival, evolved through crossover and mutation, and frozen only after three-seed finalist selection. Held-out transfer is an audit after the route is frozen.

For a capability-specific route targeting capability t , let T_t contain its source cells and let $C_t = G \setminus T_t$ contain all non-target cells. We define

$$\mu_t(S) = \frac{1}{|T_t|} \sum_{g \in T_t} \Delta_g(S), \quad w_t(S) = \max_{g \in T_t} \Delta_g(S),$$

$$\sigma_t(S) = \sqrt{\frac{1}{|T_t|} \sum_{g \in T_t} (\Delta_g(S) - \mu_t(S))^2},$$

and measure mean collateral damage as

$$\kappa_t(S) = \frac{1}{|C_t|} \sum_{g \in C_t} \Delta_g(S).$$

Evolutionary survival minimizes $q_t(S) = (\mu_t, w_t, \kappa_t, \sigma_t)$, and finalist ranking uses

$$L_t(S) = 0.45\mu_t(S) + 0.30w_t(S) + 0.15\kappa_t(S) + 0.10\sigma_t(S).$$

The collateral term discourages a route from preserving its target by destroying every other capability. It remains a soft penalty rather than a hard constraint. The worst-source and variability terms similarly prevent a good mean from concealing severe damage on one source.

3.4 Evolutionary route search

The search space contains fixed-size block subsets, so we use a seeded genetic algorithm with NSGA-II-style multi-objective survival [Deb et al., 2002]. Each route family and skip budget is searched independently. The procedure is summarized in Figure 2.

Initialization. At most half of each population is filled with prior routes constructed from the least-damaging single-block rankings and earlier combined-route searches, resized to the current value of K . The remaining positions are unique, seeded random K -subsets of the allowed vision blocks. Priors improve search efficiency but do not restrict which blocks later mutation can introduce.

Pareto survival. A candidate S_a dominates S_b when it is no worse on every component of the relevant objective vector and strictly better on at least one:

$$S_a \prec S_b \iff [\forall j, q_j(S_a) \leq q_j(S_b)] \wedge [\exists j, q_j(S_a) < q_j(S_b)].$$

Candidates are divided into nondominated fronts, which are accepted in order until half the population survives. If only part of a front fits, we prefer higher NSGA-II crowding distance, then lower scalar loss, then lexicographically smaller routes. Survivors are copied unchanged into the next generation. Before reproduction they are ordered by scalar loss and route index. For ordered survivors $R_0, \dots, R_{m-1}$, offspring attempt i pairs

$$A_i = R_{i \bmod m}, \quad B_i = R_{(5i+1) \bmod m}.$$

This deterministic schedule replaces stochastic parent selection.

Fixed- K crossover. Given two same-size parent routes A and B , the child retains every block in their intersection and fills its remaining positions from their non-shared union:

$$C = (A \cap B) \cup U, \quad U \subseteq (A \cup B) \setminus (A \cap B), \quad |U| = K - |A \cap B|.$$

The subset U is chosen by seeded SHA-256 ordering. This preserves shared parental choices while guaranteeing that the child still skips exactly K blocks.

One-swap mutation. Every crossed child is mutated exactly once by removing one included block r and adding one allowed block a outside the child:

$$M(C) = (C \setminus \{r\}) \cup \{a\}.$$

The removal and addition are again selected by seeded SHA-256 ordering. Duplicate children are rejected, and a deterministic random route is inserted only if repeated crossover-mutation attempts cannot refill the population. There is no crossover or mutation probability: every offspring receives both operations once, while elite survivors receive neither. Fitness acts through Pareto survival and deterministic survivor ordering rather than roulette-wheel or tournament selection.

3.5 Finalist selection and matched controls

We run each route-family and budget search with three fixed seeds: 20260715, 20260716, and 20260717. Qwen uses a population of 16 routes and three evaluated generations, including the initial population, for $K \in \{4, 6, 8\}$. Each seed contributes at most two Pareto finalists. After deduplication across seeds, all finalists are evaluated on a 300-example development set. At most three routes with the lowest matching scalar loss advance to the 876-example image-disjoint selection set, and the route with the lowest selection loss is frozen. Exact ties are broken by lexicographic block order.

The SmolVLM2 replication follows the same objectives, operators, three seeds, and finalist procedure, but uses a population of 12, two evaluated generations, and only $K = 4$. These leaner settings were frozen before the replication run. Routes are never automatically reused across budgets or architectures.

Every final comparison is made at matched K . The controls are: the K blocks with the lowest independent average sensitivity, the K blocks with the lowest capability-specific sensitivity, the least-damaging contiguous K -block window, and three deterministic random K -block routes. The unpruned model is reported as an accuracy reference, not as a matched-compute control. The 876-example set is used for method selection and is not described as sealed evidence because earlier discovery work touched its underlying benchmark sources. External transfer data is accessed only after routes, objectives, prompts, decoding, and scoring rules are frozen.

We quantify uncertainty with 10,000 fixed-seed paired bootstrap resamples over the same example IDs. Reported intervals are percentile 95% intervals for the accuracy difference between two routes. We treat an interval that crosses zero as inconclusive; an interval that only touches zero is reported as borderline rather than as a clear positive result.

4 Experiments

4.1 Models, datasets, and metrics

Models. The primary model is Qwen2.5-VL-3B-Instruct, which contains 32 vision blocks [Bai et al., 2025]. The replication model is SmolVLM2-2.2B-Instruct, which contains 27 vision blocks [Hugging Face, 2025]. We pin Qwen revision 66285546d2b821cf421d4f5eb2576359d3770cd3 and SmolVLM2 revision 482adb537c021c86670beed01cd58990d01e72e4. Both models use deterministic greedy decoding and batch-size-one inference. The full-model accuracies on the 876-example selection partition are 83.68% for Qwen and 82.65% for SmolVLM2.

Data. The discovery corpus contains 1,780 image-question pairs over 1,431 unique images. It combines controlled MME perception questions [Fu et al., 2023] with one larger source for each capability: OCRBench for direct text recognition [Liu et al., 2023b], TallyQA for counting [Acharya et al., 2019], POPE for object existence [Li et al., 2023], VQAv2 color questions for attributes [Goyal et al., 2017], and VSR for spatial relationships [Liu et al., 2023a]. Images are assigned as complete hash groups, producing 904 development examples and an image-disjoint 876-example selection partition. The selection partition contains 166 attribute, 181 counting, 193 object, 155 OCR, and 181 spatial questions.

The development and selection partitions are separated by image, but both source families participated in earlier single-block discovery. We therefore use the 876 examples as method-selection evidence rather than calling them a sealed test set. A separate 250-example IIIT5K scene-text subset [Mishra et al., 2012] is kept sealed until the SmolVLM2 shared and OCR-specific four-block routes are frozen. This final set tests source transfer only; it does not influence route selection, objectives, prompts, decoding, or scoring.

Metrics. The primary metric is short-answer accuracy, reported overall and within each capability. Route comparisons use paired accuracy differences in percentage points with the bootstrap intervals defined in Section 3.5. We also report mean pairwise Jaccard overlap among the three seed-level route winners, the fraction of vision and total-model parameters bypassed by identity skipping, and batch-size-one vision-encoder and end-to-end latency. Parameter counts describe inactive execution, not the size of the saved checkpoint.

4.2 Does evolutionary search beat naive pruning?

Figure 3 compares route construction methods across Qwen skip budgets. The searched shared route achieves 81.28%, 79.11%, and 75.91% accuracy when four, six, and eight blocks are skipped. Independent ranking reaches 80.25%, 78.08%, and 72.37%, so evolutionary search improves accuracy by 1.03, 1.03, and 3.54 points respectively. The advantage becomes much larger against weaker controls at aggressive budgets. With eight blocks skipped, contiguous removal reaches 40.98% and the mean of three random routes reaches 46.96%, compared with 75.91% for the searched shared route.


The same pattern appears for capability-specific construction. On Qwen, the searched capability policy exceeds independently constructed capability routes by 0.11 points with four blocks skipped, 3.77 points with six skipped, and 8.45 points with eight skipped. This increasing gap supports the central motivation: one-block scores become less reliable as more interventions must compose.

The four-block comparison replicates on SmolVLM2 (Figure 4). Its searched shared route reaches 72.49%, compared with 67.58% for independent ranking, a gain of 4.91 points with paired 95% interval [1.83, 7.99]. Its searched capability policy reaches 73.29%, compared with 64.38% for independently constructed capability routes, a gain of 8.90 points [5.82, 11.99]. The contiguous route reaches 71.58% and the random mean reaches 54.53%. Evolutionary route construction therefore improves same-budget pruning on both architectures, even though their absolute robustness differs.

4.3 Does capability-specific routing help?

Table 1 separates the search benefit from the stronger claim that capability-specific policies outperform one shared route. With four Qwen blocks skipped, the shared route and capability policy score 81.28% and 81.39%. The difference is only +0.11 points, and its interval [-1.83, 2.05] crosses zero. With eight blocks skipped, the policy gains +0.68 points [-1.71, 3.08], which is also inconclusive.

The strongest Qwen specialization result occurs with six blocks skipped. The capability policy reaches 81.28%, compared with 79.11% for the shared route, for a +2.17-point difference [0.00, 4.34]. The interval touches zero, so the overall result is borderline rather than unambiguously positive. At the capability level, OCR provides clearer evidence: its selected route gains 7.10 points [0.65, 14.19] over the shared route. Attribute gains 1.20 points, counting 1.10, spatial 2.21, and object is unchanged, but all four intervals include zero. Qwen therefore



SmolVLM2 is omitted because only the four-block study was completed.

Figure 3: Qwen selection accuracy as the number of skipped blocks increases. Every condition at a given x-axis position skips exactly the same number of vision blocks. SmolVLM2 is omitted because only its four-block study was completed.

Table 1: Matched-budget selection accuracy. A shared route uses one block set for every question; a capability-specific policy uses one same-size route per known capability label.

Model	Skipped	Full	Shared	Capability	Δ (pp)	95% CI
Qwen2.5-VL-3B	4 blocks	83.68	81.28	81.39	+0.11	[-1.83, 2.05]
Qwen2.5-VL-3B	6 blocks	83.68	79.11	81.28	+2.17	[0.00, 4.34]
Qwen2.5-VL-3B	8 blocks	83.68	75.91	76.60	+0.68	[-1.71, 3.08]
SmolVLM2-2.2B	4 blocks	82.65	72.49	73.29	+0.80	[-1.94, 3.54]

supplies evidence for useful specialization at one budget and capability, not a general rule that every capability needs a different route.

4.4 Does the result replicate across models?

SmolVLM2 replicates the value of evolutionary search but not the Qwen capability profile. Its four-block capability policy is only 0.80 points above its shared route overall, with interval [-1.94, 3.54]. That small aggregate difference hides large effects in opposite directions (Figure 5). Counting improves by 7.18 points [1.10, 13.26] and spatial reasoning improves by 9.39 points [2.76, 16.02]. OCR moves in the opposite direction, falling by 13.55 points [-21.94, -5.16]. Attribute changes by -0.60 points and object by -0.52, with both intervals crossing zero.

The sign reversal for OCR is especially important. A capability-specific OCR route helps Qwen at six skipped blocks but harms SmolVLM2 at four. Similar aggregate policy scores can therefore conceal fundamentally different behavior by capability. The result argues against interpreting one successful route as a universal capability-to-depth map. Plausible explanations include architecture-specific block interactions, source-specific shortcuts, and noise from a finite search budget; this experiment cannot isolate which explanation dominates.

Matched four-block controls: search beats naive route construction




Figure 4: Accuracy for matched four-block controls. Dashed lines indicate the corresponding full-model accuracy. Search improves route construction on both architectures, while SmolVLM2 remains more damaged by four skipped blocks overall.




Figure 5: Capability-policy accuracy minus shared-route accuracy. Stars mark paired 95% intervals that exclude zero. Qwen is shown at its strongest specialization budget, while SmolVLM2 is shown at its completed four-block budget.

4.5 Does an OCR-specific route transfer to IIIT5K?

We test the frozen SmolVLM2 routes on 250 IIIT5K word images that were never accessed during route search. The full model reaches 94.8% exact-match accuracy. The searched shared four-block route reaches 86.4%, while the OCR-specific four-block route reaches only 72.8% (Figure 6). The route selected specifically for OCR is therefore 13.6 points worse than the shared route, with paired 95% interval $[-19.2, -8.4]$.

This negative transfer rules out the strongest interpretation of the capability-specific hypothesis. The selected SmolVLM2 OCR route is not a stable OCR pathway across sources. It may have adapted to the mixture of MME and OCRBench examples used during selection, exploited architecture-specific interactions that do not support IIIT5K text, or won because several low-stability candidates had similar development losses. The experiment demonstrates the failure but cannot distinguish these mechanisms. It also shows why a capability name alone may be too coarse for routing: source, text style, image complexity, and prompt format may matter alongside the nominal task.

4.6 Efficiency and route stability

The selected block identities are often unstable across search seeds. Figure 7 reports the mean pairwise Jaccard overlap among three development winners at the four-block budget. For the shared route, overlap is 0.095 on Qwen and 0.159 on SmolVLM2. Qwen’s attribute and OCR searches are more stable at 0.733 and 0.556, but




Figure 6: Post-freeze source transfer on IIIT5K. Both compressed conditions skip four SmolVLM2 vision blocks. Routes and all evaluation choices were frozen before any IIIT5K prediction was produced.

its spatial winners have zero overlap. Every SmolVLM2 route family remains at or below 0.222. Different block sets can therefore produce similar objective values, which further weakens any claim that the search recovers one canonical pathway for a capability.

Identity skipping produces a meaningful reduction in executed vision depth but a modest reduction in whole-model parameters (Figure 8). Four Qwen blocks account for 11.79% of vision-tower parameters and 2.10% of total model parameters. Four SmolVLM2 blocks account for 14.76% of vision parameters and 2.71% of total parameters. These values describe parameters bypassed during execution. The original modules remain in the loaded checkpoint, so our implementation does not yet produce a physically smaller model file.

For the searched SmolVLM2 shared four-block route, repeated batch-size-one measurements show an 8.60% vision-encoder speedup and a 4.19% end-to-end speedup on the same RTX 4090 VM. The smaller end-to-end gain is expected because image preprocessing and autoregressive language generation are unchanged. We do not report a final Qwen accuracy-latency curve: its committed latency audit predates the final evolved routes and does not provide matched measurements across the four-, six-, and eight-block winners.

Together, the results support four conclusions. First, block interactions make route construction combinatorial, and evolutionary search is consistently better than naive composition. Second, a safe pruning budget is architecture-dependent: SmolVLM2 loses 10.16 points from its full model even under its searched shared four-block route, while Qwen loses 2.40. Third, aggregate accuracy can hide severe capability-level gains and losses. Fourth, capability labels alone are not reliable routing variables. A learned router conditioned on question, source, image complexity, and hardware cost is a promising extension, but it is not evaluated here.





Figure 7: Mean pairwise Jaccard overlap among three seed-level four-block development winners. Low overlap indicates that the frozen route is one strong solution among several distinct candidates, not a uniquely identified block map.

K4 removes vision depth, not most total model parameters



Unlocked same-VM fallback; not fixed-clock or edge-device evidence.

Figure 8: Four-block parameter fractions for both models and measured SmolVLM2 latency. The latency measurement is an unlocked same-VM comparison on one RTX 4090, not fixed-clock or edge-device evidence.

5 Limitations

This study covers two model architectures, five visual capability labels, and one constructed mixture of short-answer benchmarks. The evolutionary budget is finite, so every selected route is the best found under the frozen search rather than a globally optimal subset. Dataset identity also remains partly confounded with capability. Although MME provides a second controlled source within each bucket and the objective balances capability-source cells, OCRBench, TallyQA, POPE, VQAv2, and VSR still contribute different image and answer distributions. Public benchmark examples may also have appeared in model pretraining data.

The 876-example selection partition is image-disjoint from route development, but it is not a pristine held-out benchmark because its source families participated in earlier single-block screening. Only the 250-example IIIT5K audit is a genuinely fresh post-freeze source-transfer test, and it evaluates one capability on one model. The paired bootstrap intervals quantify uncertainty over the observed examples. They do not capture uncertainty across model training runs, benchmark families, search hyperparameters, or alternative capability taxonomies.

Our capability-specific policy assumes an oracle capability label. We do not train a per-question router, measure routing mistakes, or include routing overhead. Identity skipping updates no weights, so the study does not test recovery through fine-tuning, distillation, or low-rank adaptation. It also evaluates deterministic short-answer correctness rather than long-form generation, grounding, calibration, or safety behavior.

Finally, our efficiency evidence is limited. The implementation bypasses blocks at runtime but retains their parameters in the checkpoint. It does not produce a deployable compact checkpoint, prune the language decoder, quantize weights, or combine depth skipping with visual-token reduction. Latency was measured at batch size one on one RTX 4090 VPS. The SmolVLM2 result is an unlocked same-VM comparison because fixed GPU clocks were unavailable, and no mobile or edge hardware, memory bandwidth, energy consumption, or complete Qwen final-route latency series was measured.

6 Related Work

Structured depth and width reduction. ShortGPT measures residual-block influence and removes low-influence language-model layers without retraining [Men et al., 2024]. SliceGPT exploits computational invariance to delete rows and columns from transformer weight matrices, reducing hidden width rather than depth [Ashkboos et al., 2024]. These methods establish that structured compression can map to dense hardware operations, but neither asks whether combinations of vision-encoder blocks should vary by visual capability.

Short-LVLM directly studies training-free layer pruning in large vision-language models [Ma et al., 2025]. It uses important vision-language tokens to improve layer localization and applies a low-rank subspace compensation to reduce the feature gap after pruning. Its objective is a generally compressed model, whereas our primary empirical question compares one shared route with a capability-specific policy and tests whether independently safe vision blocks compose at matched skip budgets. The methods are complementary: compensation could be applied after our route search, while our results caution that localization quality alone may not predict multi-block compatibility.

Visual-token reduction. DynamicViT learns input-dependent token sparsification within vision transformers [Rao et al., 2021], while Token Merging combines similar tokens without retraining [Bolya et al., 2023]. SparseVLM and VScan reduce visual-token computation in large vision-language models [Zhang et al., 2025b,a], and FlowCut analyzes information flow to identify redundant visual computation [Tong et al., 2025]. These approaches primarily reduce sequence width or token processing. Our intervention reduces executed vision-encoder depth, making route search potentially compatible with token reduction rather than a replacement for it.

Conditional VLM pruning. FlashVLM uses visual attention to guide inference-time layer skipping during VLM decoding [Sarkar et al., 2026]. Domain-aware pruning work studies how decoder-layer rankings change across domains [Khaki et al., 2026]. This work shares the broad idea that redundancy can depend on the input distribution, but differs in intervention target and evaluation. We search fixed-cardinality subsets in the vision tower, compare shared and named-capability routes at identical budgets, and explicitly test cross-architecture behavior, source transfer, and seed stability. Our findings do not imply that task-aware VLM efficiency is otherwise unexplored; they show that capability-conditioned vision-depth routes are less stable than aggregate method-selection accuracy suggests.

7 Conclusion

We asked whether combinations of vision-encoder blocks can be skipped more intelligently than independent layer ranking, and whether different visual capabilities benefit from different routes. Across Qwen2.5-VL-3B and SmolVLM2-2.2B, the first answer is consistently yes: source-balanced evolutionary search constructs stronger same-budget routes than independent, contiguous, or random selection. The method is especially valuable as the skip budget grows and block interactions make one-block rankings unreliable.

The second answer is conditional. Qwen shows a useful six-block capability policy, including a clear OCR gain, while SmolVLM2 shows large counting and spatial gains alongside a larger OCR loss. The sealed IIIT5K test confirms that the SmolVLM2 OCR-specific route does not transfer even within the capability it was designed to preserve. We therefore do not find a stable capability-to-layer map. The practical result is narrower but reproducible: combinatorial search is a useful tool for vision-depth pruning, while reliable dynamic routing will require richer per-input signals, stronger multi-source validation, and hardware-aware optimization.

References

- Manoj Acharya, Kushal Kafle, and Christopher Kanan. Tallyqa: Answering complex counting questions. In *AAAI Conference on Artificial Intelligence*, 2019. URL <https://arxiv.org/abs/1810.12440>.
- Saleh Ashkboos, Maximilian L. Croci, Marcelo Gennari do Nascimento, Torsten Hoefer, and James Hensman. Slicept: Compress large language models by deleting rows and columns. In *International Conference on Learning Representations*, 2024. URL <https://arxiv.org/abs/2401.15024>.
- Shuai Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Sibong Song, Kai Dang, Peng Wang, Shijie Wang, Jun Tang, et al. Qwen2.5-vl technical report. *arXiv preprint arXiv:2502.13923*, 2025. URL <https://arxiv.org/abs/2502.13923>.
- Daniel Bolya, Cheng-Yang Fu, Xiaoliang Dai, Peizhao Zhang, Christoph Feichtenhofer, and Judy Hoffman. Token merging: Your vit but faster. In *International Conference on Learning Representations*, 2023. URL <https://arxiv.org/abs/2210.09461>.
- Kalyanmoy Deb, Amrit Pratap, Sameer Agarwal, and T. Meyarivan. A fast and elitist multiobjective genetic algorithm: Nsga-ii. *IEEE Transactions on Evolutionary Computation*, 6(2):182–197, 2002. doi: 10.1109/4235.996017.
- Chaoyou Fu, Peixian Chen, Yunhang Shen, Yulei Qin, Mengdan Zhang, Xu Lin, Jinrui Yang, Xiawu Zheng, Ke Li, Xing Sun, Yunsheng Wu, and Rongrong Ji. Mme: A comprehensive evaluation benchmark for multimodal large language models. *arXiv preprint arXiv:2306.13394*, 2023. URL <https://arxiv.org/abs/2306.13394>.
- Yash Goyal, Tejas Khot, Douglas Summers-Stay, Dhruv Batra, and Devi Parikh. Making the v in vqa matter: Elevating the role of image understanding in visual question answering. In *IEEE Conference on Computer Vision and Pattern Recognition*, 2017. URL <https://arxiv.org/abs/1612.00837>.
- Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, et al. Training compute-optimal large language models. *arXiv preprint arXiv:2203.15556*, 2022. URL <https://arxiv.org/abs/2203.15556>.
- Hugging Face. Smolvlm2-2.2b-instruct model card. Hugging Face model repository, 2025. URL <https://huggingface.co/HuggingFaceTB/SmolVLM2-2.2B-Instruct>. Revision 482adb537c021c86670beed01cd58990d01e72e4 used in this work.
- Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models. *arXiv preprint arXiv:2001.08361*, 2020. URL <https://arxiv.org/abs/2001.08361>.
- Saeed Khaki, Nima Safaei, and Kamal Ginotra. Understanding pruning regimes in vision-language models through domain-aware layer selection. *arXiv preprint arXiv:2603.20275*, 2026. URL <https://arxiv.org/abs/2603.20275>.
- Yifan Li, Yifan Du, Kun Zhou, Jinpeng Wang, Wayne Xin Zhao, and Ji-Rong Wen. Evaluating object hallucination in large vision-language models. *arXiv preprint arXiv:2305.10355*, 2023. URL <https://arxiv.org/abs/2305.10355>.
- Fangyu Liu, Guy Emerson, and Nigel Collier. Visual spatial reasoning. In *Transactions of the Association for Computational Linguistics*, 2023a. URL <https://arxiv.org/abs/2205.00363>.

- Yuliang Liu, Zhang Li, Biao Li, Chunyuan Yu, Xucheng Huang, Yuming Liu, Lianwen Jin, and Xiang Bai. On the hidden mystery of ocr in large multimodal models. *arXiv preprint arXiv:2305.07895*, 2023b. URL <https://arxiv.org/abs/2305.07895>.
- Ji Ma, Wei Suo, Peng Wang, and Yanning Zhang. Short-lvlm: Compressing and accelerating large vision-language models by pruning redundant layers. In *Proceedings of the 33rd ACM International Conference on Multimedia*, 2025. doi: 10.1145/3746027.3755122. URL <https://arxiv.org/abs/2507.23362>.
- Xin Men, Mingyu Xu, Qingyu Zhang, Bingning Wang, Hongyu Lin, Yaojie Lu, Xianpei Han, and Weipeng Chen. Shortgpt: Layers in large language models are more redundant than you expect. *arXiv preprint arXiv:2403.03853*, 2024. URL <https://arxiv.org/abs/2403.03853>.
- Anand Mishra, Karteek Alahari, and C. V. Jawahar. Scene text recognition using higher order language priors. In *British Machine Vision Conference*, 2012. doi: 10.5244/C.26.127.
- Yongming Rao, Wenliang Zhao, Benlin Liu, Jiwen Lu, Jie Zhou, and Cho-Jui Hsieh. Dynamicvit: Efficient vision transformers with dynamic token sparsification. In *Advances in Neural Information Processing Systems*, 2021. URL <https://arxiv.org/abs/2106.02034>.
- Sreetama Sarkar, Saket Gururkar, Yuexi Zhang, Monami Banerjee, and Ashwin Chandra. Flashvlm: Exploiting layer redundancy via visual attention for efficient vision-language inference. *OpenReview preprint*, 2026. URL <https://openreview.net/forum?id=7U9pt9yUQF>.
- Jintao Tong, Wenwei Jin, Pengda Qin, Anqi Li, Yixiong Zou, Yuhong Li, Yuhua Li, and Ruixuan Li. Flowcut: Rethinking redundancy via information flow for efficient vision-language models. In *Advances in Neural Information Processing Systems*, 2025. URL <https://openreview.net/forum?id=M6zQNbCaLl>.
- Ce Zhang, Kaixin Ma, Tianqing Fang, Wenhao Yu, Hongming Zhang, Zhisong Zhang, Yaqi Xie, Katia P. Sycara, Haitao Mi, and Dong Yu. Vscan: A two-stage visual token reduction framework for accelerating large vision-language models. In *Findings of the Association for Computational Linguistics: ACL 2025*, 2025a. URL <https://arxiv.org/abs/2505.22654>.
- Yuan Zhang, Chun-Kai Fan, Junpeng Ma, Wenzhao Zheng, Tao Huang, Kuan Cheng, Denis Gudovskiy, Tomoyuki Okuno, Yohei Nakata, Kurt Keutzer, et al. Sparsevlm: Visual token sparsification for efficient vision-language model inference. In *International Conference on Machine Learning*, 2025b. URL <https://arxiv.org/abs/2410.04417>.